

- Bandwidth
 - Ingress / Egress
- Any data eviction policy?
- Any spike in traffic?
- Storage
 - distributed file storage systems: hdfs, glusterfs
 - any data eviction policy?
 - Non volatile storage
- Cache
 - any cache eviction policy? LRU?
- Load Balancer -> SSL termination
- Data replication
- Monitoring
- Security
- Any Machine Learning pattern is usable?
- Lambda Architecture
- CDN
- Datacenter as a service
- 429 - too many requests
- Race condition
- Rate limiter: fixed, rolling
- UUID : Universally unique ID
- Session persistance
- Stale data
- Domino effect - we want the replicas to Have independent failure probability to Have uncorrelated failure

- !!! Data - communication - computation
- What about utilization?
- Quorum for read / write?
- For replication; replicated state machine (read this, write that); state transfer (transfer all database). state transfer more robust. RSM has a problem with parallelism.
- What is output rule for system design, mentioned in VMware paper?
- What is Split brain disaster?
- What is test and set server?
- Automated canary analyses?

- Autoscaling? for stateless and stateful machines? Immutable machine images?
 - Consistent Naming?
 - ELB Config?
 - Healthcheck?
 - Squeeze testing
 - Staged, red/black deployments
 - timeouts / retries / fallbacks
 - circuit breakers, fallbacks, chaos
 - multi-region failover
 - failure-driven design
 - Write back, write ahead caches???
 - Serializability : concurrency control
 - Linearizability : consistency
 - Concurrency control : Pessimistic -> two phase locking, optimistic -> ?
 - Web sockets - full duplex - allows communication in both direction - unlike half-duplex, allow this to happen simultaneously.
-
- We begin with the problem statement, gather requirements, and iterate through designs that become increasingly sophisticated until we reach a viable solution. Ultimately, we arrive at a system that defends against many failure modes and satisfies both the initial requirements and additional details that emerged as we iterated.
 - I will try to split my system into services that each one do one thing and do it well.
 - Gives different parts of my system isolation from each other.
 - Makes it easier to troubleshoot problems.
 - Makes it easier for the system to evolve over time.
 - Simplifies performance analysis. No need to think about different traffic mixes.
 - I will try to keep the state in specific components and keep the rest of the system stateless.
 - Stateless components easy to manage, because they are trivial to scale horizontally.
 - Flow
 - First design for single DC, then multiple DC
 - Assumption, Constraints, SLOs

- Authentication and Authorization already taken care of? So all users are already authenticated and trusted clients?
- API
- Data (Critical / Non-critical)
- Do some calculations
 - Disk (1 TB SSD / HDD)
 - RAM (128 GB)
 - Bandwidth (10 Gbps -> 1 GBps)
 - CPU (32 cores)
 - Datacenters are connected via 100Gbps?
 - Assume the probability of a machine failing some number of times?
 - 0 -> 40 %
 - 1 -> 40 %
 - 2 -> 15 %
 - 3 -> 4 %
 - >3 -> 1%
- Iterate & Refine (If I am unhappy with any aspect)
 - Is it feasible and does that work correctly?
 - Is it reliable?
 - What is the scariest thing that can happen to my system? And how my system defend against it?
 - Loss of a datacenter
 - Loss of important state / data
 - Load spikes
 - Efficient use of hardware? Utilization?
 - Are there places where I can improve by adding caching, indexing, lookup filters like bloom filters.
- Monitoring
 - Detect outages (and alert) - Auto scalers, commodity computing
 - Give indications of growth and system capacity to meet future needs
 - Assist in troubleshooting outages and problems
- Consistency
 - FastPaxos : Optimized consensus algorithm that uses a stable leader process to improve performance. Performance for FP is

one round-trip between the leader and quorum of closest replicas to commit an update or a consistent read.

- Sometimes keys are not evenly distributed. It is always best to shard based on a hash of the key.
 - $M \geq 100 * N$
 - $\text{shard} = \text{hash}(\text{key}) \% M$
 - we have a map of shard -> server
- Sharded and replicated datastores should have automatic mechanisms to cross-check state with other replicas and load lost states from the peers. Maybe after a downtime or when new replicas are added. This needs to be rate-limited somehow avoid thundering-herd problems where many replicas are out of sync.

External consistency

At most one semantics

- Technical (expertise in features, systems, processes; technical discussions with the engineers)
- Ownership (scope of projects, ability to deliver projects, coordination)
- Business Insight (understanding team delivery; setting priorities for business value, vision)
- Continuous Improvement (mentorship and coaching, championing changes, well-being of the team)
- Leadership (communication across organisation, building communities, helping recruiting processes))

- Mainframe hash çalışması
- Google Hashcode hub world championship
- Slack group - data structures & algorithms
- son 5 dakika için ek sorular
- LinkedIn Common Questions - evernote